




APPLYING GANs IN THE MEDICAL INDUSTRY

ARTURO POLANCO LOZANO
Universidad Surcolombiana



CONTENT

1. Generative Modelling
2. Generative Adversarial Networks
3. Data
4. Generating Medical Images with GANs
5. Segmentation of Medical Images with GANs




Generative VS Discriminative

Discriminative modeling estimates $p(y|\mathbf{x})$ —the probability of a label y given observation \mathbf{x} .

Generative modeling estimates $p(\mathbf{x})$ —the probability of observing observation \mathbf{x} .

If the dataset is labeled, we can also build a generative model that estimates the distribution $p(\mathbf{x}|y)$.

Generative VS Discriminative

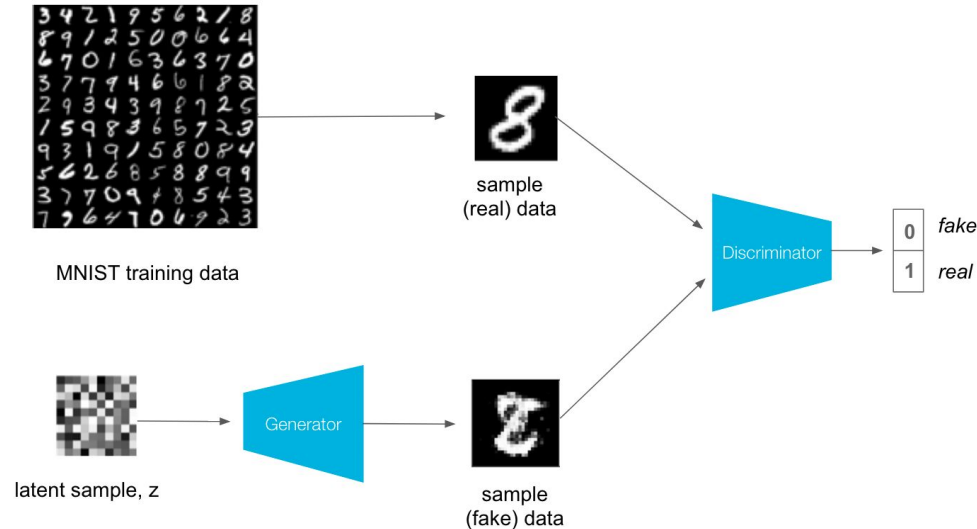


CONTENT

1. Generative Modelling
2. Generative Adversarial Networks
3. Data
4. Generating Medical Images with GANs
5. Segmentation of Medical Images with GANs

Generative Adversarial Networks

A GAN is a generative model that is trained using two neural network models. One model is called the “*generator*” or “*generative network*” model that learns to generate new plausible samples. The other model is called the “*discriminator*” or “*discriminative network*” and learns to differentiate generated examples from real examples.





Generator

Discriminator



Real Money



Counterfeiter prints fake money. It is labelled as fake for police training. Sometimes, the counterfeiter attempts to fool the police by labelling the fake money as real.



Fake Money



The police are trained to spot real from fake money. Sometimes, the police give feedback to the counterfeiter why the money is fake.

GAN Lab

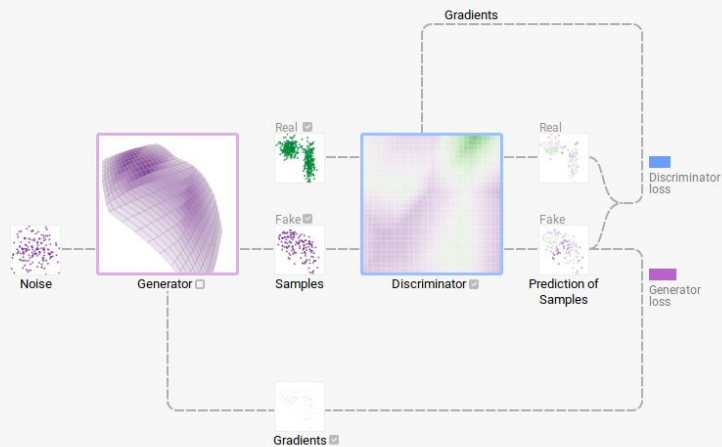
Data Distribution

☒ Use pre-trained model


Epoch

001,931

MODEL OVERVIEW GRAPH




LAYERED DISTRIBUTIONS

Each dot is a 2D data sample: real samples fake samples.Background colors of grid cells represent discriminator's classifications. Samples in green regions are likely to be real; those in purple regions likely fake.Manifold represents generator's transformation results from noise space. Opacity encodes density: darker purple means more samples in smaller area.Pink lines from fake samples represent gradients for generator.

This sample needs to move upper right to decrease generator's loss.

METRICS

■ Discriminator's Loss■ Generator's Loss■ KL Divergence (by grid)■ JS Divergence (by grid)

This GAN Lab webpage records anonymous click information (e.g., buttons in the GAN Lab interface clicked). It will be used for studying the usage behavior of users and improving the usability of GAN Lab. If you are not interested in contributing to the study, [click here](#) to remove your data.

<https://poloclub.github.io/ganlab/>

This Website Generates AI Portraits of People Who Don't Exist

FEB 19, 2019

MICHAEL ZHANG

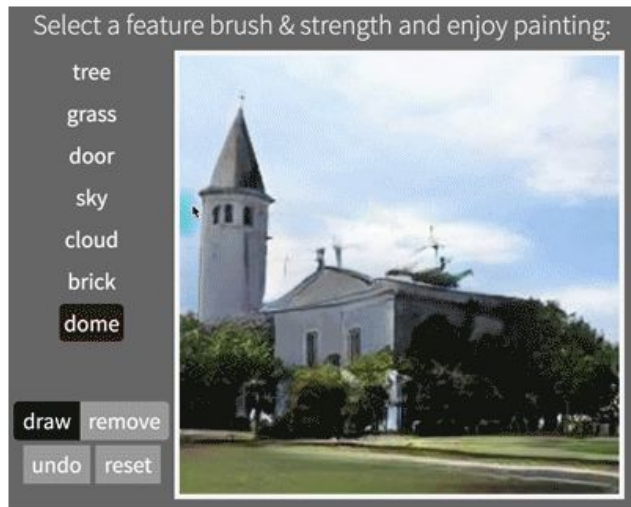
Share 253

Tweet

21 COMMENTS



<https://thispersondoesnotexist.com/>



The [#GANpaint app](https://gandissect.csail.mit.edu/) works by directly activating and deactivating sets of neurons in a deep network trained to generate images. Each button on the left ("door", "brick", etc) corresponds to a set of 20 neurons. The app demonstrates that, by learning to draw, the network also learns about objects such as trees and doors and rooftops. By switching neurons directly, you can observe the structure of the visual world that the network has learned to model. ([Try it here.](#))

<https://gandissect.csail.mit.edu/>



Image-to-Image Demo

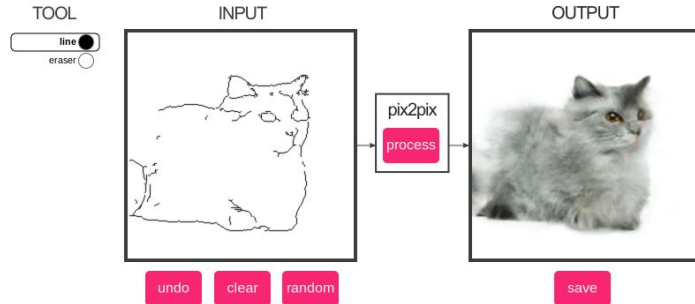
Interactive Image Translation with pix2pix-tensorflow

Written by *Christopher Hesse* — February 19th 2017

Recently, I made a *Tensorflow* port of *pix2pix* by *Isola et al.*, covered in the article *Image-to-Image Translation in Tensorflow*. I've taken a few pre-trained models and made an interactive web thing for trying them out. Chrome is recommended.

The *pix2pix* model works by training on pairs of images such as building facade labels to building facades, and then attempts to generate the corresponding output image from any input image you give it. The idea is straight from the *pix2pix* paper, which is a good read.

edges2cats



<https://affinelayer.com/pixsrv/>

Applications of GANs in Images

- Generate Examples for Image Datasets
- Generate Photographs of Human Faces
- Generate Realistic Photographs
- Generate Cartoon Characters
- Image-to-Image Translation
- Text-to-Image Translation
- Semantic-Image-to-Photo Translation
- Face Frontal View Generation
- Generate New Human Poses
- Photos to Emojis
- Photograph Editing
- Face Aging
- Photo Blending
- Super Resolution
- Photo Inpainting
- Clothing Translation
- Video Prediction
- 3D Object Generation

Applications of GANs in Images Conditional AN



Applications of GANs in Images DC GAN




(a)

(b)



(c)

(d)



Applications of GANs in Images (StackGan)

this small bird has a pink breast and crown, and black primaries and secondaries.



this magnificent fellow is almost all black with a red crest, and white cheek patch.



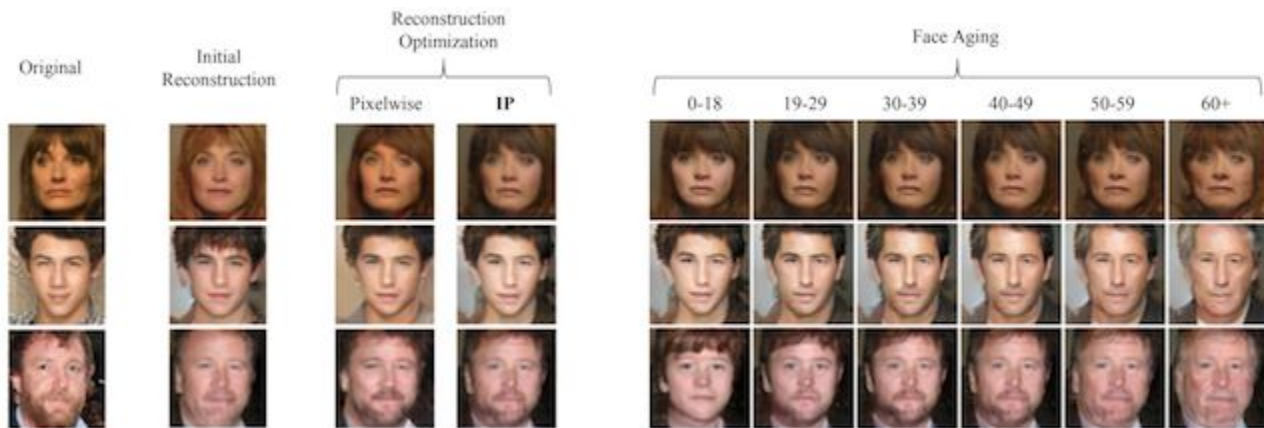
the flower has petals that are bright pinkish purple with white stigma



this white and yellow flower have thin white petals and a round yellow stamen



Applications of GANs in Images Conditional GAN




Applications of GANs in Images (SRGAN)



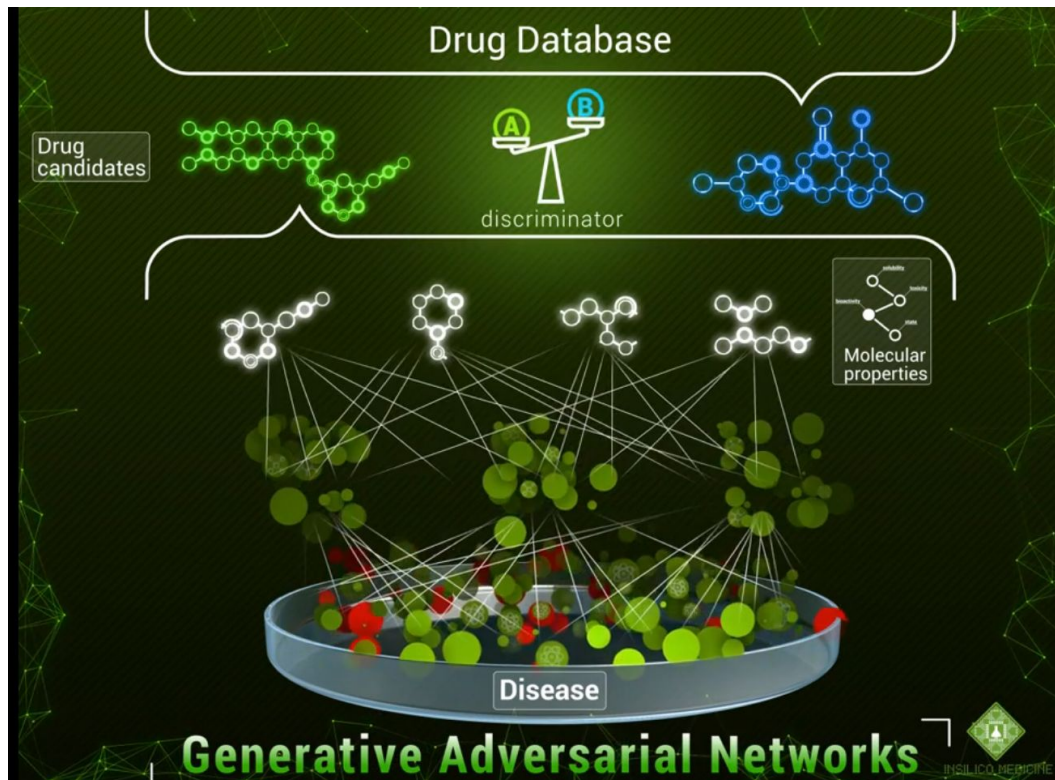
Applications of GANs in Images



Applications of GANs in Images



Other Applications of GANs



Other Applications of GANs



Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis

Luke de Oliveira^a, Michela Paganini^{a,b}, and Benjamin Nachman^a

^a*Lawrence Berkeley National Laboratory, 1 Cyclotron Rd, Berkeley, CA, 94720, USA*

^b*Department of Physics, Yale University, New Haven, CT 06520, USA*

Other Applications of GANs



Visit The GAN Zoo on Github !

The GAN Zoo



Every week, new GAN papers are coming out and it's hard to keep track of them all, not to mention the incredibly creative ways in which researchers are naming these GANs! So, here's a list of what started as a fun activity compiling all named GANs!

CONTENT

1. Generative Modelling
2. Generative Adversarial Networks
3. Data
4. Generating Medical Images with GANs
5. Segmentation of Medical Images with GANs

Multimodal Brain Tumor Segmentation Challenge 2018



• [Scope](#) • [Relevance](#) • [Tasks](#) • [Data](#) • [Evaluation](#) • [Participation Summary](#) • [Data Request](#) • [Previous BraTS](#) • [People](#) •

The background of the slide is a grid of brain MRI slices. The top row features four axial slices showing various brain structures and lesions. The bottom row shows the top portions of several more axial slices. A semi-transparent gray rectangle is centered over the middle of the grid, containing the title text.

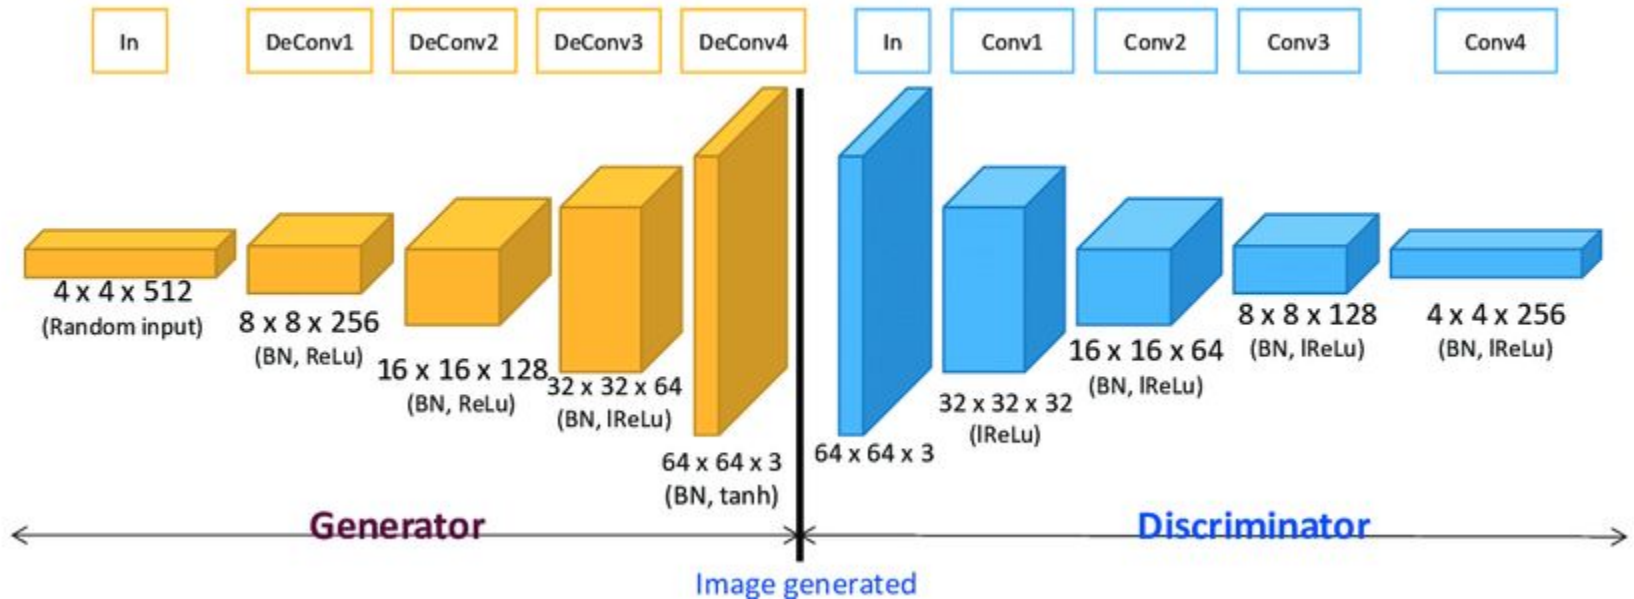
ISLES CHALLENGE 2018

ISCHEMIC STROKE
LESION
SEGMENTATION

CONTENT

1. Generative Modelling
2. Generative Adversarial Networks
3. Data
4. *Generating Medical Images with GANs*
5. Segmentation of Medical Images with GANs

Deep Convolutional GAN



Discriminator

```
class Discriminator(tf.keras.Model):
    def __init__(self):
        super(Discriminator, self).__init__()
        self.conv1 = tf.keras.layers.Conv2D(64, (4, 4), strides=(2, 2), padding='same')
        self.conv2 = tf.keras.layers.Conv2D(64 * 2, (4, 4), strides=(2, 2), padding='same')
        self.batchnorm2 = tf.keras.layers.BatchNormalization()
        self.conv3 = tf.keras.layers.Conv2D(64 * 4, (4, 4), strides=(2, 2), padding='same')
        self.batchnorm3 = tf.keras.layers.BatchNormalization()
        self.conv4 = tf.keras.layers.Conv2D(64 * 8, (4, 4), strides=(2, 2), padding='same')
        self.batchnorm4 = tf.keras.layers.BatchNormalization()
        self.conv5 = tf.keras.layers.Conv2D(1, (4, 4), strides=(1, 1), padding='valid')

        self.dropout = tf.keras.layers.Dropout(0.3)
        self.flatten = tf.keras.layers.Flatten()
        self.fcl = tf.keras.layers.Dense(1)

    def call(self, x, training=True):
        x = tf.nn.leaky_relu(self.conv1(x))

        x = self.conv2(x)
        x = self.batchnorm2(x, training=training)
        x = tf.nn.leaky_relu(x)

        x = self.conv3(x)
        x = self.batchnorm3(x, training=training)
        x = tf.nn.leaky_relu(x)

        x = self.conv4(x)
        x = self.batchnorm4(x, training=training)
        x = tf.nn.leaky_relu(x)

        x = self.conv5(x)
        x = self.fcl(x)
        return x
```

Generator

```
class Generator(tf.keras.Model):
    def __init__(self):
        super(Generator, self).__init__()
        self.fc1 = tf.keras.layers.Dense(8*8*64, use_bias=False)
        self.batchnorm1 = tf.keras.layers.BatchNormalization()

        self.conv1 = tf.keras.layers.Conv2DTranspose(64 * 8, (4, 4), strides=(1, 1), padding='same', use_bias=False)
        self.batchnorm2 = tf.keras.layers.BatchNormalization()

        self.conv2 = tf.keras.layers.Conv2DTranspose(64 * 4, (4, 4), strides=(2, 2), padding='same', use_bias=False)
        self.batchnorm3 = tf.keras.layers.BatchNormalization()

        self.conv3 = tf.keras.layers.Conv2DTranspose(64 * 2, (4, 4), strides=(2, 2), padding='same', use_bias=False)
        self.batchnorm4 = tf.keras.layers.BatchNormalization()

        self.conv4 = tf.keras.layers.Conv2DTranspose(64 * 1, (4, 4), strides=(2, 2), padding='same', use_bias=False)
        self.batchnorm5 = tf.keras.layers.BatchNormalization()

        self.conv5 = tf.keras.layers.Conv2DTranspose(1, (4, 4), strides=(1, 1), padding='same', use_bias=False)

    def call(self, x, training=True):
        x = self.fc1(x)
        x = self.batchnorm1(x, training=training)
        x = tf.nn.relu(x)

        x = tf.reshape(x, shape=(-1, 8, 8, 64))

        x = self.conv1(x)
        x = self.batchnorm2(x, training=training)
        x = tf.nn.relu(x)

        x = self.conv2(x)
        x = self.batchnorm3(x, training=training)
        x = tf.nn.relu(x)

        x = self.conv3(x)
        x = self.batchnorm4(x, training=training)
        x = tf.nn.relu(x)


        x = self.conv4(x)
        x = self.batchnorm5(x, training=training)
        x = tf.nn.relu(x)

        x = tf.nn.tanh(self.conv5(x))
        return x
```




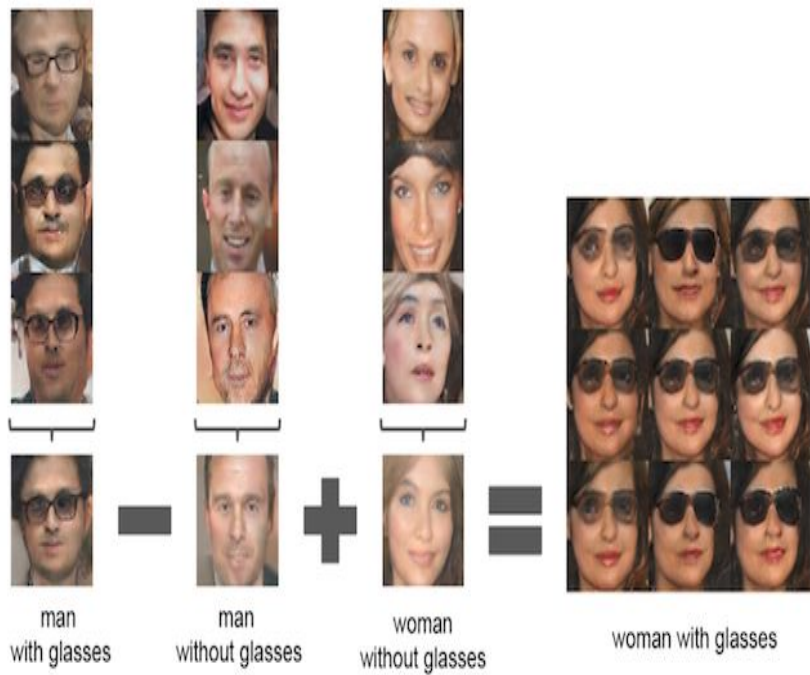
Discriminator and Generator Loss

```
def discriminator_loss(real_output, generated_output):  
    # [1,1,...,1] with real output since it is true and we want  
    # our generated examples to look like it  
    real_loss = tf.losses.sigmoid_cross_entropy(multi_class_labels=tf.ones_like(real_output), logits=real_output)  
  
    # [0,0,...,0] with generated images since they are fake  
    generated_loss = tf.losses.sigmoid_cross_entropy(multi_class_labels=tf.zeros_like(generated_output),  
    logits=generated_output)  
  
    total_loss = real_loss + generated_loss  
  
    return total_loss  
  
def generator_loss(generated_output):  
    return tf.losses.sigmoid_cross_entropy(tf.ones_like(generated_output), generated_output)
```

Output



Explore Latent Space



Goals

- In order to use larger deep learning models is required hughes amounts of data, GANs can be used for data augmentation instead of other methods that involve image cropping, translation, etc.
- Some pathologic findings are rare, GANs can be used for generating images with specific features.
- Medical data needs to remain private, using GANs in order to create synthetic images helps to solve data anonymization problem.

CONTENT

1. Generative Modelling
2. Generative Adversarial Networks
3. Data
4. Generating Medical Images with GANs
5. Segmentation of Medical Images with GANs

Im2Im Translation with Conditional AN








Figure 3: Two choices for the architecture of the generator. The “U-Net” [50] is an encoder-decoder with skip connections between mirrored layers in the encoder and decoder stacks.

Output



Output



Goals

- Comparison between the amount of gray matter, white matter and cerebrospinal fluid are vital for doctors in order to perform diagnostics.
- The diagnostic of ischemic stroke is time sensitive, a delay in the treatment could signify in permanent damage or death of the patient.

GAN Hacks

Architecture guidelines for stable Deep Convolutional GANs

- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

GAN Hacks

[soumith / ganhacks](#)

[Watch](#) 332[Star](#) 7.3k[Fork](#) 1.2k

[Code](#)[Issues 43](#)[Pull requests 3](#)[Projects 0](#)[Wiki](#)[Security](#)[Insights](#)

starter from "How to Train a GAN?" at NIPS2016

5 commits


1 branch

0 releases


2 contributors


Branch: master [New pull request](#)


[Create new file](#)[Upload files](#)[Find file](#)[Clone or download](#)

 **soumith** Merge pull request #4 from raminia/dropout [...](#)

Latest commit d99bbf3 on Dec 16, 2016

 [images](#) first commit 3 years ago

 [README.md](#) Added a dropout trick based on <https://arxiv.org/pdf/1611.07004v1.pdf> 3 years ago

 [README.md](#)

How to Train a GAN? Tips and tricks to make GANs work

While research in Generative Adversarial Networks (GANs) continues to improve the fundamental stability of these models, we use a bunch of tricks to train them and make them stable day to day.

Here are a summary of some of the tricks.

Thanks :-)