




# What to Wear?



Outfit recommendation using  
computer vision and machine  
learning





# Motivation





# Motivation

---

- You may spend a lot of time deciding what to wear.
- It can affect your self esteem.
- It impacts the way others see you.
- There are multiple etiquettes.
- Depends on where you are.
- Changes on time.

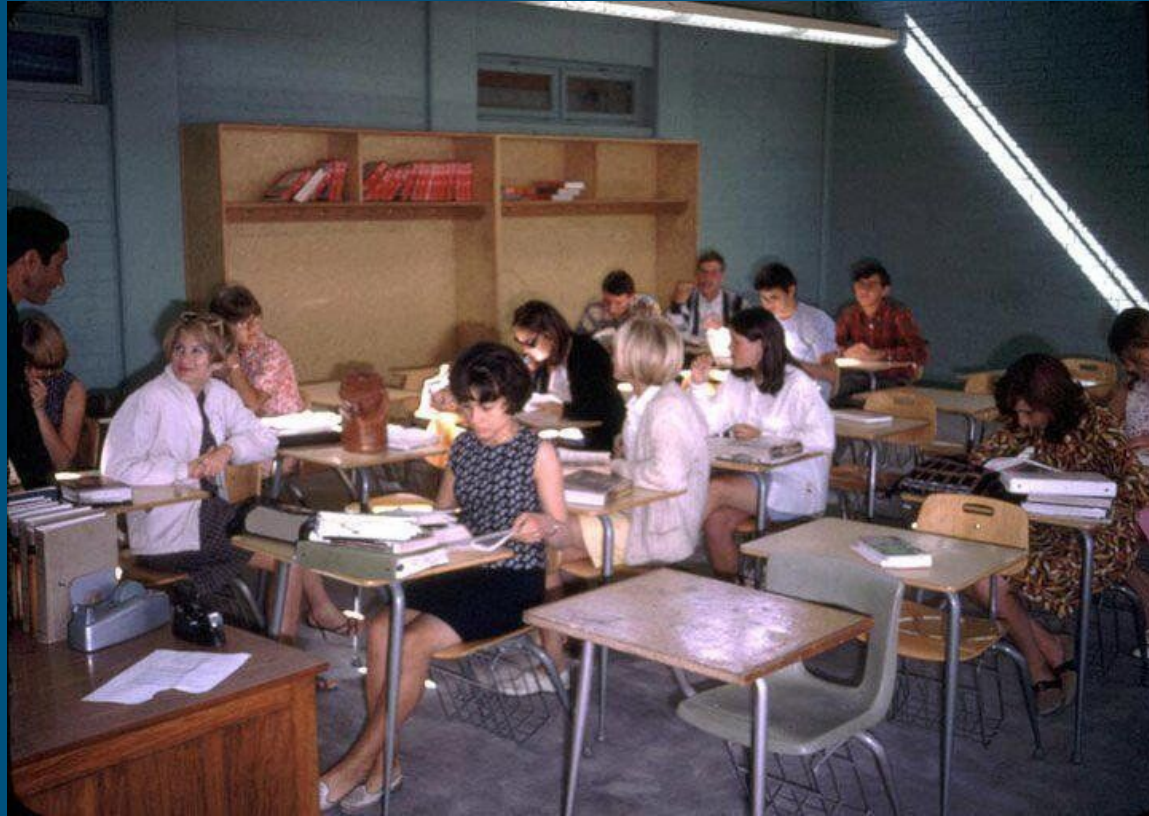
# Different places, different “rules”

---



# Different times, different “rules”

---



# Who am I?

---

- Computer and Systems Engineer, UTP.
- Student of Master in CS, AI Branch, UTP.
- Software Developer for 6 years.
- Married to a Fashion Designer.
- My master thesis was basically this.

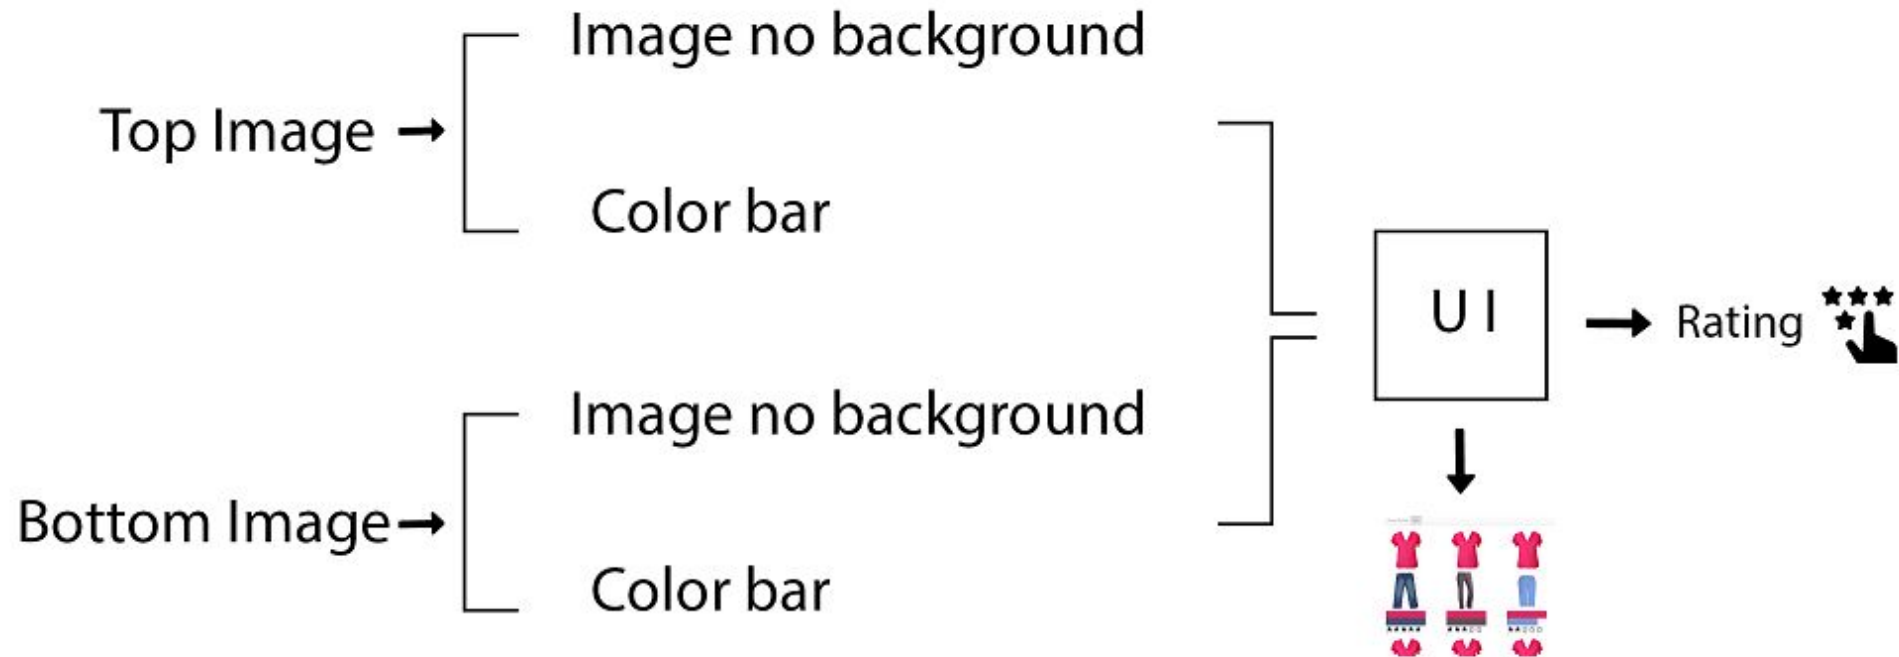
# Challenges

---

1. Build a DB for this very specific purpose.
2. Segmentate the clothes.
3. Extract features that are relevant.
4. Assemble those features.
5. Train a model to predict rating on 2 new images.
6. Put it all together in a usable way.



# Building a DB



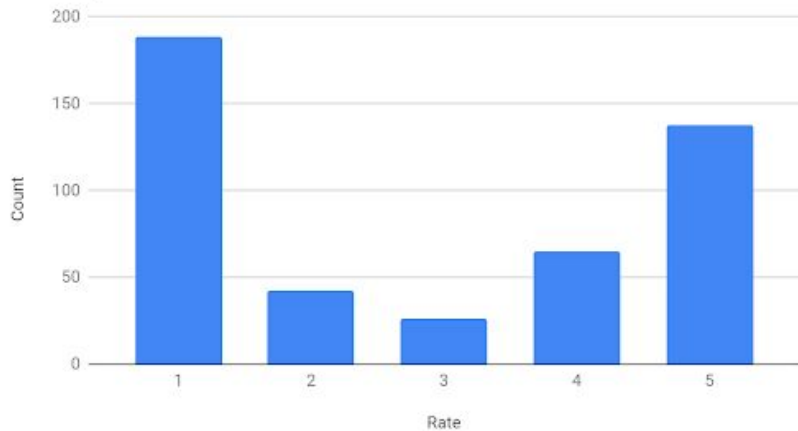
# Building a DB: Capturing rating

---

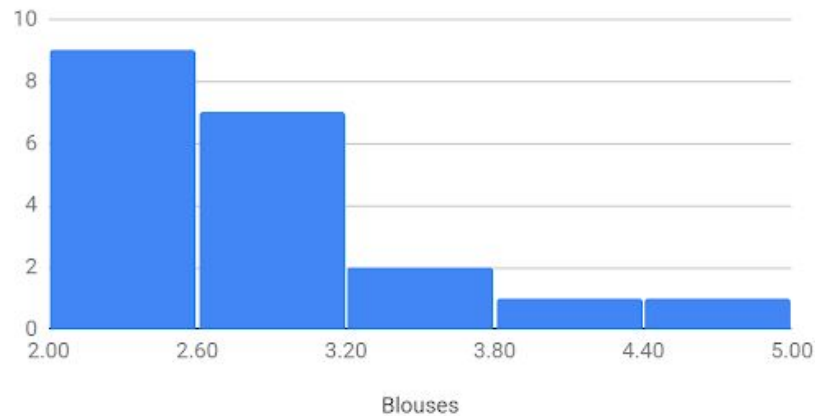


# Building a DB, First attempt: terrible clothes

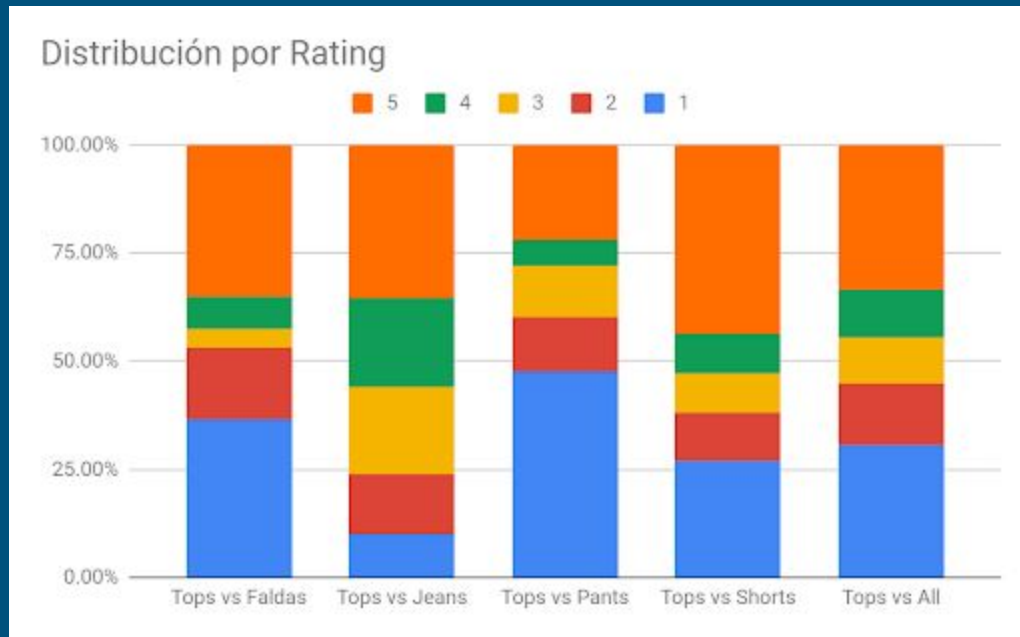
Count per Rate



Histogram of Blouses



# Building a DB, Second attempt: cool clothes.

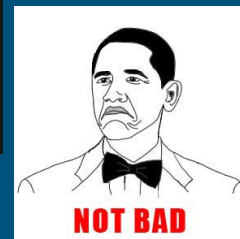


# Segmentation: GrabCut

---

1. User selects a bounding box.
2. Estimates color distribution of target object and background with Gaussian Mixture Model.
3. Builds a Markov Random Field con energy function.
4. Runs a GraphCut based optimization.
5. Repeats until convergence.
6. Can be corrected by user by pointing misclassified regions and re running.

# Segmentation: GrabCut

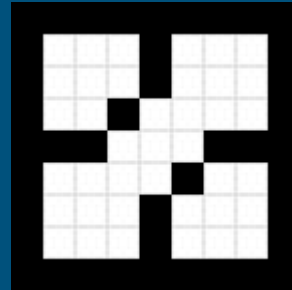


# Segmentation: Floodfill

---

1. If target-color is equal to replacement-color, return.
2. Else if the color of node is not equal to target-color, return.
3. Else Set the color of node to replacement-color.
4. Perform Flood-fill to each direction.
5. Done.

Our replacement color in RGBA will have  $\alpha=0$



# Segmentation: Floodfill

---





# Segmentation: Comparison

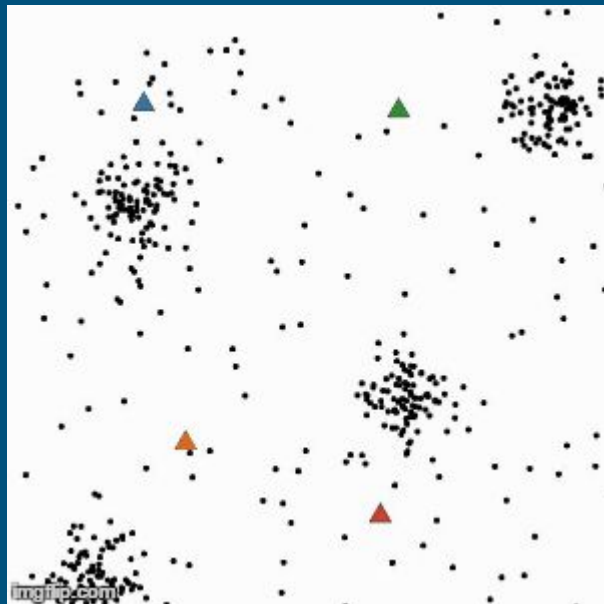
---

- 83% of cases FloodFill is as good or better than GrabCut.
- GrabCut was as good 38% of the time and better 13%.
- For this case, Floodfill is clearly better! (plus, it's a 100 times faster).
- GrabCut could be more precise though, especially when backgrounds are not plain.

# Main Colors: K-Means

---

- Find clusters
- Minimize variance inside the group
- Maximize variance between groups
- You need to know K
- My approach for K: heuristic.
  - Start with 3 colors.
  - If 2 of the found colors are too similar (less than 15% difference), run with lower K.
  - If 2 of the found colors are similar (less than 20% difference) and one of them represents less than 5% of the pixels, run with lower K.
  - If K equals 1, you may want to stop now.



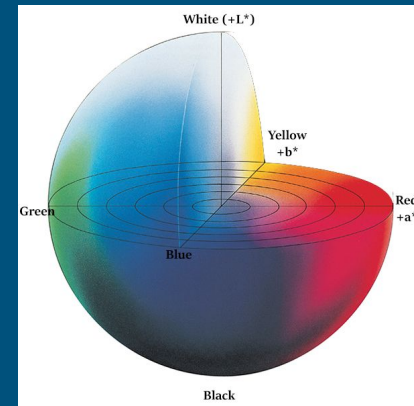
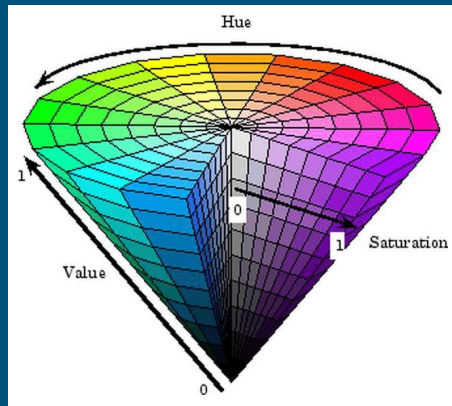
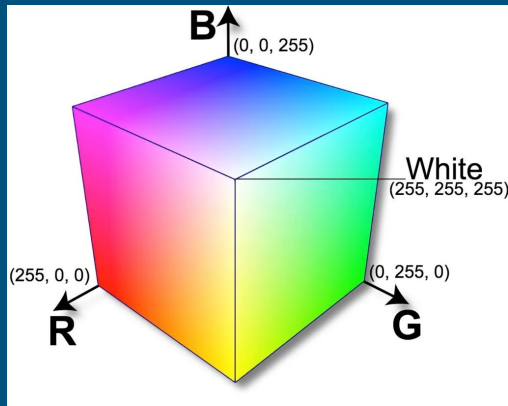
# Main Colors: K-Means

---



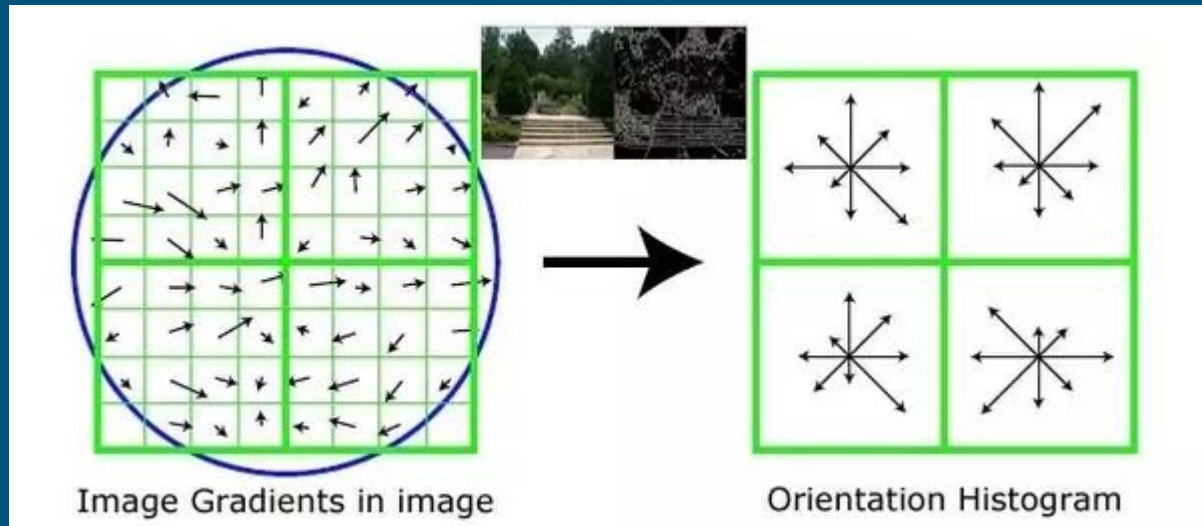
# Features: Color Histograms

- Color spaces:
  - RGB: red, green, blue.
  - HSV: hue, saturation, value.
  - Lab: lightness, green to red, blue to yellow.
- Bins: 8, for each color component



# Features: HOG, Histogram of Oriented Gradients

Counts occurrences of gradient orientation in localized portions of an image.



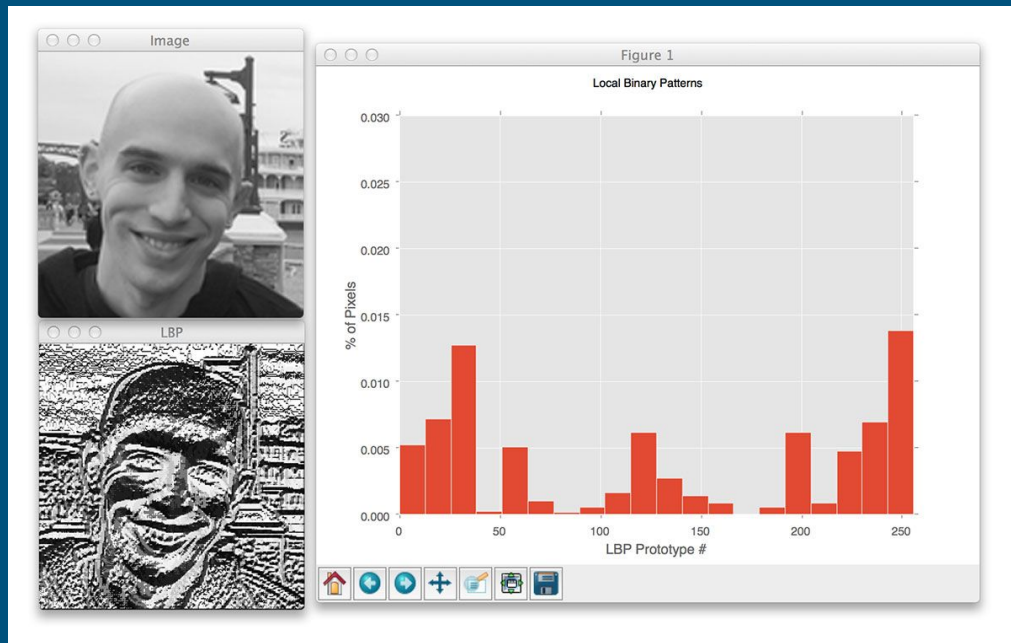
# Features: LBP, Local Binary Patterns

Texture descriptor which builds local representations by comparing each pixel with its surrounding neighborhood of pixels.

Check <https://www.pyimagesearch.com/>!

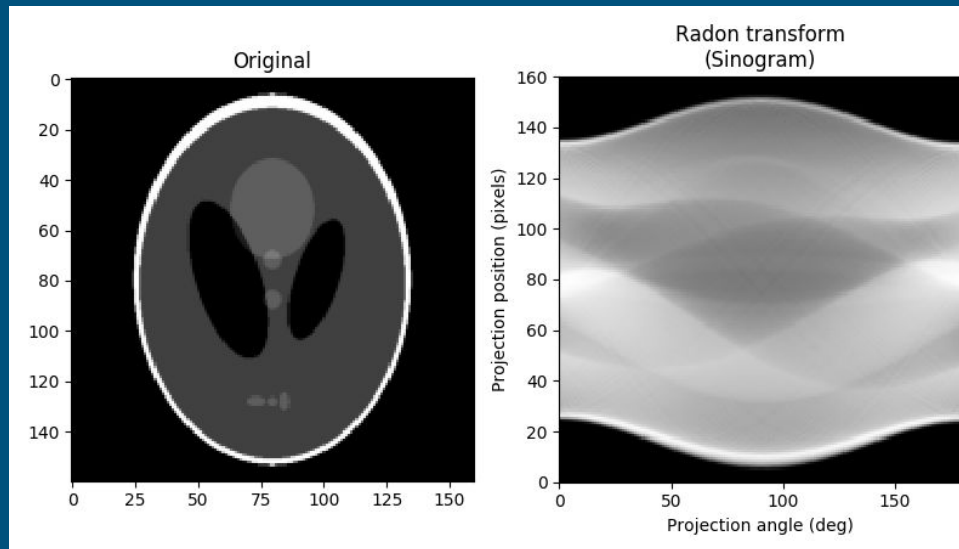
You may also want to check:

<https://www.learnopencv.com/>

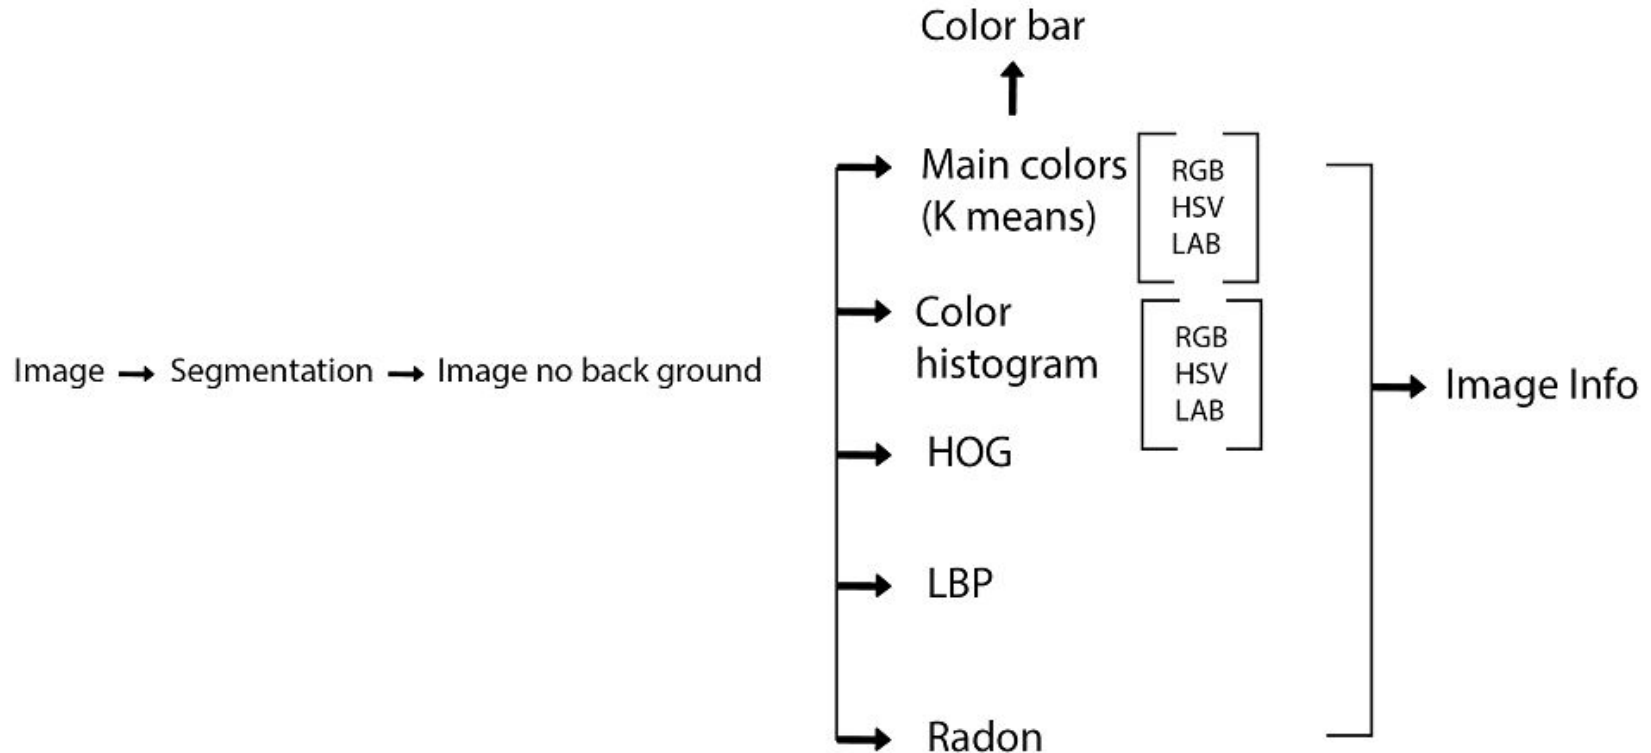


# Features: Radon Signature

The Radon transform is the integral transform which takes a function  $f$  defined on the plane to a function  $Rf$  defined on the (two-dimensional) space of lines in the plane, whose value at a particular line is equal to the line integral of the function over that line



# Features: Overview





# Assembly: Feature options

---

Several choices:

- Use main colors extended to 1, 2, 3 or 5 colors, or use color histogram [5].
- Use RGB, HSV, LAB, just H or all of them [5].
- Use LBP with radius 2, 5, 10 or simply not use it [4].
- Use HOG or not [2].
- Use Radon transform or not [2].
- Use dimension reduction: PCA, ICA or none [3].

1200 combinations so far.

# Training: Models

---

**Neural Networks:** Set of artificial neurons connected with weights and an activation function.

**Bayes Regression:** Statistical analysis using Bayesian Inference.

**SVM with RBF kernel:** Take points to a new space where they can be separated by a Hyperplane. RFB is radial basis function.

**Random Forest:** Train a set of decision trees with different samples, then use the mode or the average to predict new cases.

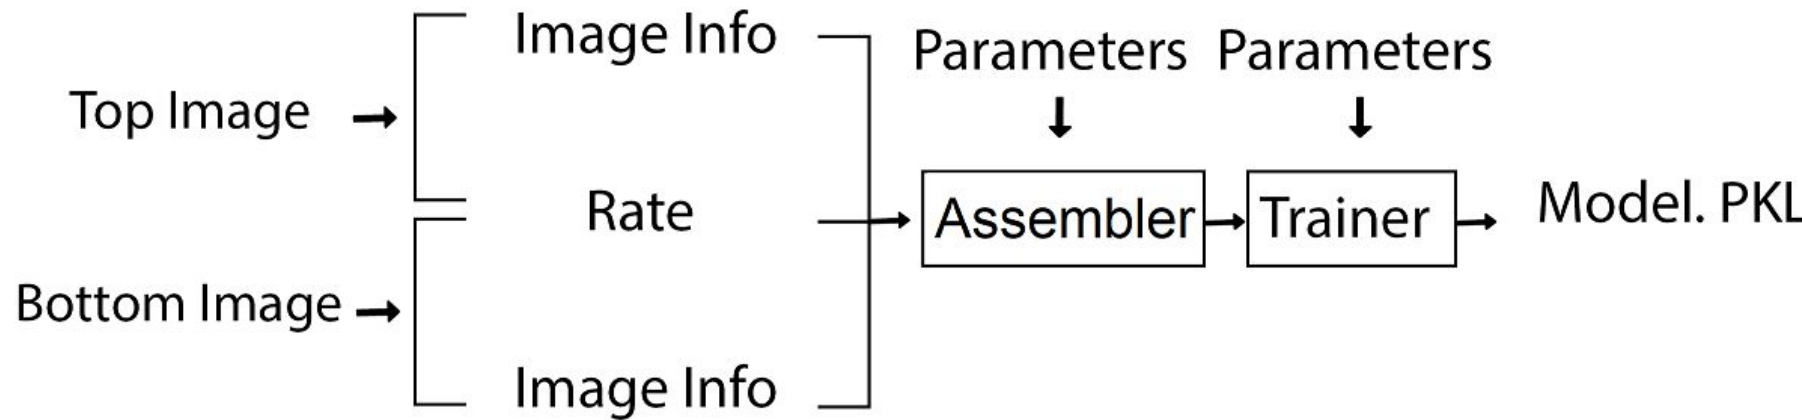
# Training: Model options

---

- SVM: C can be  $10^1$ ,  $10^2$ ,  $10^3$ ,  $10^4$ . Gamma  $10^{-1}$ ,  $10^{-2}$ . [8]
- Neural Networks. Vary the number of hidden layers. [4]
- Random Forest. Num of trees: 50, 100, 200. Max depth: 5, 6, 7, 8. [12]
- Bayes Regression. [1]

25 unique configurations.

# Training: Overview



# Training: All options

---

I tested each independent category (tops vs pants, tops vs ...) and all together. [5]

25 model options \* 1200 assembly options \* 5 sets = 150k.

If you do cross validation, WHICH YOU SHOULD, you get (with 5 folds):

# 750k unique trainings



# Training: GridSearchCV

---

- Performs exhaustive search.
- Receives a param grid.
- Has flexible scoring.
- Allows cv with Stratified KFold or custom splitting.
- Finds the best parameters.
- Gives you a dictionary with cv\_results.
- You'll need more than that.



```

def train_and_validate(
    tops, bottoms, num_folds, names_and_model_functions,
    color_spaces_options, n_jobs, colors_per_image_options,
    save_best_model=False, run_for_each=True, run_for_all=True,
    **kwargs,
):
    if run_for_each:
        1 for top in tops:
            2 for bottom in bottoms:
                3 for color_spaces in color_spaces_options:
                    4 for colors_per_image in colors_per_image_options:
                        X, y, dimensions = get_x_and_y_from_options([top], [bottom], color_spaces, colors_per_image, **kwargs)
                        kwargs['dimensions'] = dimensions

                        5 for model_name, model_fun, model_params in names_and_model_functions:
                            log_current(top, bottom, color_spaces, model_name, colors_per_image, **kwargs)
                            grid = train_configuration(model_name, model_fun, model_params, X, y, num_folds, n_jobs)
                            full_set_score = validate_model(grid.best_estimator_, X, y)
                            results = get_results_from_grid(
                                grid=grid,
                                top=top,
                                bottom=bottom,
                                model_name=model_name,
                                color_spaces=color_spaces,
                                num_folds=num_folds,
                                colors_per_image=colors_per_image,
                                **kwargs,
                            )
                            results['full_set_score'] = full_set_score
                            yield results

                        if save_best_model:
                            filename = get_model_filename(top, bottom, model_name, colors_per_image)
                            joblib.dump(grid.best_estimator_, filename)

```





# Training: GridSearchCV not enough?

Try HungaBunga!

“Brute Force all scikit-learn models and all scikit-learn parameters with fit predict.”

<https://github.com/ypeleg/HungaBunga>

`HungaBungaClassifier.fit(X, y)`



`HungaBungaRegressor.fit(X, y)`

# Training: Results

Bottom	Color Spaces	Colors per Img	Model	LBP	HOG	Radon	Red.	Mean Test Score	Std Test Score	Best Params
Shorts	RGB	5	Forest	No	Yes	No	PCA	0.881	0.028	Max Depth: 7, N Estimators: 50
Shorts	HSV	2	Forest	No	Yes	No	PCA	0.849	0.092	Max Depth: 8, N Estimators: 50
Jeans	RGBHSV	2	Neural Net.	Yes (R 5)	Yes	Yes	None	0.813	0.056	Hidden Layers: (50, 50, 50), Max Iter: 5000
Jeans	*	2	Neural Net.	No	Yes	No	None	0.812	0.027	Hidden Layers: (50, 50, 50), Max Iter: 5000
Skirts	RGB	5	Forest	No	No	Yes	PCA	0.801	0.017	Max Depth: 7, N Estimators: 100
Skirts	H	Histogram	Forest	No	Yes	No	PCA	0.778	0.026	Max Depth: 8, N Estimators: 50
Pants	*	3	SVM(Rbf)	No	No	Yes	ICA	0.784	0.023	C: 1000, Gamma: 0.1
Pants	RGB	Histogram	SVM(Rbf)	Yes (R 10)	Yes	Yes	None	0.784	0.022	C: 100, Gamma: 0.01
All	LAB	2	Forest	No	No	No	None	0.765	0.042	Max Depth: 8, N Estimators: 200
All	*	Histogram	Forest	Yes (R 10)	Yes	Yes	PCA	0.763	0.026	Max Depth: 8, N Estimators: 100

# Training: Results using Random Forest Only

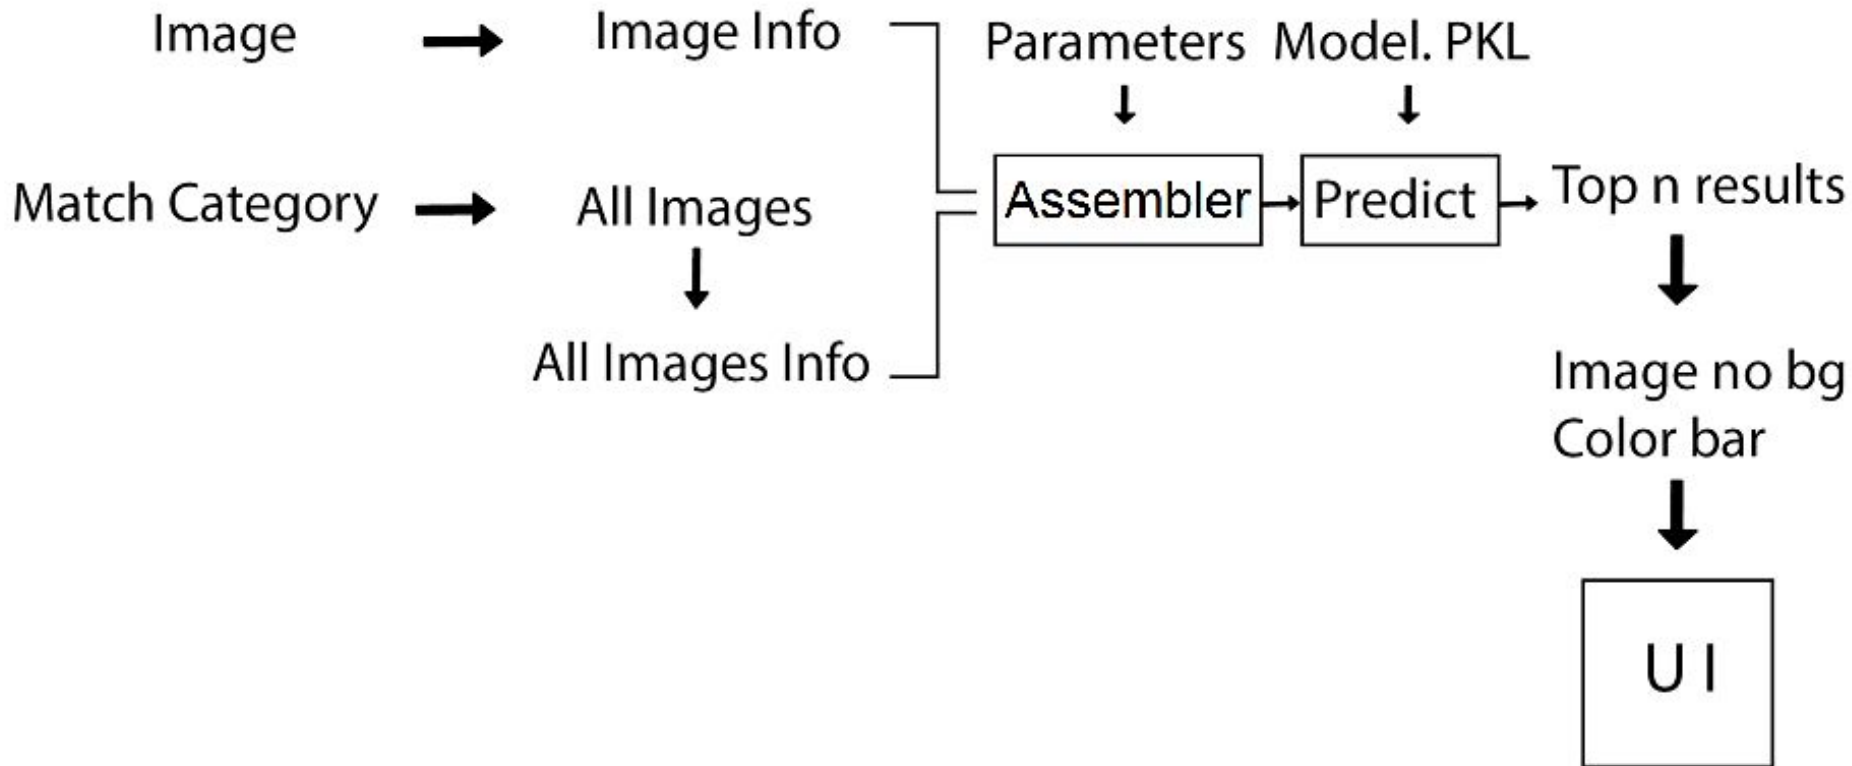
---

For random forest the best parameters in general were:

- Use 2 colors per image or histogram (cheaper).
- Color space: LAB
- No LBP, no HOG, no Radon.
- No dimension reduction.

In average we get  $74.1\% \pm 5\%$  accuracy among the 5 sets.

# Putting it all together: Overview



# Putting it all together

---



# Tools used from sklearn

---

- **GridSearchCV:** `sklearn.model_selection.GridSearchCV`
- **KMeans:** `sklearn.cluster.KMeans`
- **SVM:** `sklearn.svm`
- **Random Forest:** `sklearn.ensemble.RandomForestRegressor`
- **Bayes Regression:** `sklearn.linear_model.BayesianRidge`
- **Neural Networks:** `sklearn.neural_network.MLPRegressor`
- **PCA:** `sklearn.decomposition.PCA`
- **ICA:** `sklearn.decomposition.FastICA`

# Tools used from skimage and opencv

---

- **GrabCut:** `cv2.grabCut`
- **FloodFill:** `cv2.floodFill`
- **LBP:** `skimage.feature.local_binary_pattern`
- **Radon:** `skimage.transform.radon`
- **HOG:** `cv2.HOGDescriptor`
- **Convert color spaces:** `cv2.cvtColor`

# Learnings

---

- Analyse your data before training. It might be biased.
- Use tools as GridSearchCV.
- Define clearly your validation function, and test it.
- Use different models.
- Don't try to do every combination, do some screening first.
- Have some basic understanding of what you're using.



# Ideas for Future Work

---

1. Handle more than 2 clothes, for instance include also shoes.
2. Use GrabCut with user feedback.
3. Try other models: Gaussian Mixture, Deep Learning...
4. Try other features: ???

# Questions?

---

A word cloud featuring the phrase "Thank You" in numerous languages and colors. The central and largest text is "thank you" in red. Other prominent words include "danke" (blue), "gracias" (green), "mercí" (orange), "teşekkür ederim" (pink), "shukriya" (purple), "arigatō" (purple), "dank je" (green), "dziękuję" (pink), "obrigado" (green), "bedankt" (yellow), "spasibo" (red), "raahmat" (red), "ngiyabonga" (red), "shukra jayla" (blue), "tapadh leat" (orange), "mochchakkeram" (blue), "maith agat" (purple), "sukriya" (purple), "kop khun krap" (green), "terima kasih" (yellow), "grazie" (blue), "takk" (green), "dakujem" (orange), "merci" (orange), "sagolun" (blue), "najis tuke" (blue), "kam sah hamnida" (blue), "didi madloba" (blue), "mesí" (blue), "sagolun" (blue), "najis tuke" (blue), "kam sah hamnida" (blue), "didi madloba" (blue), "mesí" (blue), "sagolun" (blue), "najis tuke" (blue), "kam sah hamnida" (blue), "didi madloba" (blue), "mesí" (blue). The words are arranged in a circular pattern around the central "thank you".